

Abalone Age Prediction

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Load the dataset

```
data=pd.read_csv("/content/sample_data/abalone.csv")
data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight \
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

```
Age=1.5+data.Rings
data["Age"]=Age
data=data.rename(columns = {'Whole weight':'Whole_weight','Shucked weight': 'Shucked_weight','Viscera weight': 'Viscera_weight',
                             'Shell weight': 'Shell_weight'})
data=data.drop(columns=["Rings"],axis=1)
data.head()
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight \
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	


```
0.1140
4    I    0.330      0.255    0.080      0.2050      0.0895
0.0395
```

```
Shell_weight  Age
0          0.150 16.5
1          0.070  8.5
2          0.210 10.5
3          0.155 11.5
4          0.055  8.5
```

```
#Univariate Analysis
```

```
sns.displot(data["Age"])
```

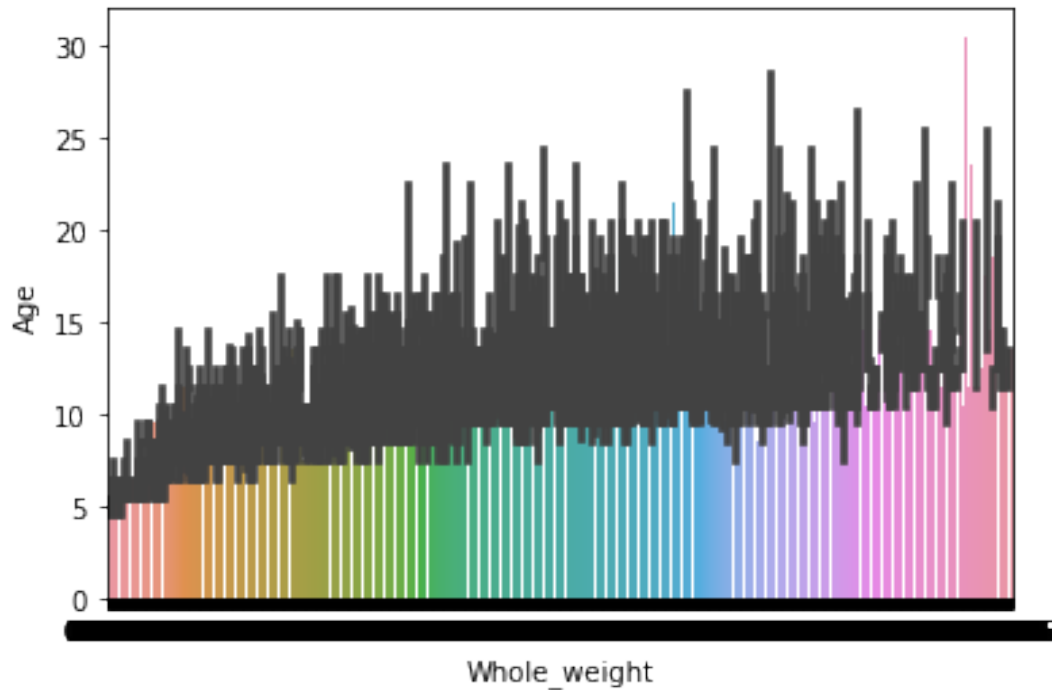
```
<seaborn.axisgrid.FacetGrid at 0x7fd9225619d0>
```



```
#Bi-Variate Analysis
```


```
sns.barplot(x=data.Whole_weight,y=data.Age)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd91fbc3950>
```



```
sns.scatterplot(x=data.Whole_weight,y=data.Height)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd91f1651d0>
```



```
#Multi-variate Analysis
```

```
sns.pairplot(data=data[["Length","Diameter","Height","Whole_weight","Shucked_weight","Viscera_weight","Shell_weight","Age",]])
```

<seaborn.axisgrid.PairGrid at 0x7fd91fbe93d0>



#Descriptive statistics

```
data.describe(include='all')
```

	Sex	Length	Diameter	Height	Whole_weight	\
count	4177	4177.000000	4177.000000	4177.000000	4177.000000	
unique	3	NaN	NaN	NaN	NaN	
top	M	NaN	NaN	NaN	NaN	
freq	1528	NaN	NaN	NaN	NaN	
mean	NaN	0.523992	0.407881	0.139516	0.828742	
std	NaN	0.120093	0.099240	0.041827	0.490389	
min	NaN	0.075000	0.055000	0.000000	0.002000	
25%	NaN	0.450000	0.350000	0.115000	0.441500	
50%	NaN	0.545000	0.425000	0.140000	0.799500	
75%	NaN	0.615000	0.480000	0.165000	1.153000	

```
max      NaN      0.815000      0.650000      1.130000      2.825500
```

```
count      Shucked_weight  Viscera_weight  Shell_weight  Age
unique      NaN            NaN            NaN            NaN
top          NaN            NaN            NaN            NaN
freq        NaN            NaN            NaN            NaN
mean        0.359367       0.180594       0.238831       11.433684
std         0.221963       0.109614       0.139203       3.224169
min         0.001000       0.000500       0.001500       2.500000
25%         0.186000       0.093500       0.130000       9.500000
50%         0.336000       0.171000       0.234000      10.500000
75%         0.502000       0.253000       0.329000      12.500000
max         1.488000       0.760000       1.005000      30.500000
```


```
data.corr()
```

```
          Length  Diameter    Height  Whole_weight
Shucked_weight \
Length          1.000000  0.986812  0.827554      0.925261
0.897914
Diameter        0.986812  1.000000  0.833684      0.925452
0.893162
Height          0.827554  0.833684  1.000000      0.819221
0.774972
Whole_weight    0.925261  0.925452  0.819221      1.000000
0.969405
Shucked_weight 0.897914  0.893162  0.774972      0.969405
1.000000
Viscera_weight 0.903018  0.899724  0.798319      0.966375
0.931961
Shell_weight    0.897706  0.905330  0.817338      0.955355
0.882617
Age            0.556720  0.574660  0.557467      0.540390
0.420884
```

```
          Viscera_weight  Shell_weight  Age
Length          0.903018      0.897706  0.556720
Diameter        0.899724      0.905330  0.574660
Height          0.798319      0.817338  0.557467
Whole_weight    0.966375      0.955355  0.540390
Shucked_weight 0.931961      0.882617  0.420884
Viscera_weight 1.000000      0.907656  0.503819
Shell_weight    0.907656      1.000000  0.627574
Age            0.503819      0.627574  1.000000
```

```
sns.heatmap(data.corr(),annot=True)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fd9180453d0>
```



#Handle The Missing values

```
data.isnull().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole_weight 0
Shucked_weight 0
Viscera_weight 0
Shell_weight 0
Age          0
dtype: int64
```

#Find the outliers

```
outliers=data.quantile(q=(0.25,0.75))
outliers
```

```
      Length  Diameter  Height  Whole_weight  Shucked_weight
Viscera_weight \
0.25    0.450     0.35   0.115         0.4415           0.186
0.0935
0.75    0.615     0.48   0.165         1.1530           0.502
```

0.2530

	Shell_weight	Age
0.25	0.130	9.5
0.75	0.329	12.5


```
a = data.Age.quantile(0.25)
b = data.Age.quantile(0.75)
c = b - a
lower_limit = a - 1.5 * c
data.median(numeric_only=True)
```

Length	0.5450
Diameter	0.4250
Height	0.1400
Whole_weight	0.7995
Shucked_weight	0.3360
Viscera_weight	0.1710
Shell_weight	0.2340
Age	10.5000

dtype: float64

```
data['Age'] = np.where(data['Age'] < lower_limit, 7, data['Age'])
sns.boxplot(x=data.Age)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fd916270e50>



#Categorical columns and encoding

```
data.head()
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight \
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

	Shell_weight	Age
0	0.150	16.5
1	0.070	8.5
2	0.210	10.5
3	0.155	11.5
4	0.055	8.5

```
from sklearn.preprocessing import LabelEncoder
labenc = LabelEncoder()
data.Sex = labenc.fit_transform(data.Sex)
data.head()
```

	Sex	Length	Diameter	Height	Whole_weight	Shucked_weight	\
0	2	0.455	0.365	0.095	0.5140	0.2245	
1	2	0.350	0.265	0.090	0.2255	0.0995	
2	0	0.530	0.420	0.135	0.6770	0.2565	
3	2	0.440	0.365	0.125	0.5160	0.2155	
4	1	0.330	0.255	0.080	0.2050	0.0895	

	Viscera_weight	Shell_weight	Age
0	0.1010	0.150	16.5
1	0.0485	0.070	8.5
2	0.1415	0.210	10.5
3	0.1140	0.155	11.5
4	0.0395	0.055	8.5

#Split the data into dependent and independent variables.

```
x=data.drop(columns=["Sex"],axis=1)
y = data["Sex"]
print(x.head())
print(y.head())
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	Viscera_weight \
0	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	0.350	0.265	0.090	0.2255	0.0995	


```

0.0485
2    0.530    0.420    0.135    0.6770    0.2565
0.1415
3    0.440    0.365    0.125    0.5160    0.2155
0.1140
4    0.330    0.255    0.080    0.2050    0.0895
0.0395

```

```

    Shell_weight  Age
0         0.150  16.5
1         0.070   8.5
2         0.210  10.5
3         0.155  11.5
4         0.055   8.5

```

```

0    2
1    2
2    0
3    2
4    1

```

Name: Sex, dtype: int64

#Scale the independent variables

```

from sklearn.preprocessing import scale
X_Scaled = pd.DataFrame(scale(x), columns=x.columns)
X_Scaled.head()

```

```

    Length  Diameter    Height  Whole_weight  Shucked_weight
Viscera_weight \
0 -0.574558 -0.432149 -1.064424    -0.641898    -0.607685    -
0.726212
1 -1.448986 -1.439929 -1.183978    -1.230277    -1.170910    -
1.205221
2  0.050033  0.122130 -0.107991    -0.309469    -0.463500    -
0.356690
3 -0.699476 -0.432149 -0.347099    -0.637819    -0.648238    -
0.607600
4 -1.615544 -1.540707 -1.423087    -1.272086    -1.215968    -
1.287337

```

```

    Shell_weight    Age
0    -0.638217  1.577830
1    -1.212987 -0.919022
2    -0.207139 -0.294809
3    -0.602294  0.017298
4    -1.320757 -0.919022

```

#Split the data into training and testing

```

from sklearn.model_selection import train_test_split
X_Train, X_Test, Y_Train, Y_Test = train_test_split(X_Scaled, y,

```

```
test_size=0.2, random_state=0)
```

```
X_Train.shape,X_Test.shape
```

```
((3341, 8), (836, 8))
```

```
Y_Train.shape,Y_Test.shape
```

```
((3341,), (836,))
```

```
X_Train.head()
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	\
3141	-2.864726	-2.750043	-1.423087	-1.622870	-1.553902	
3521	-2.573250	-2.598876	-2.020857	-1.606554	-1.551650	
883	1.132658	1.230689	0.728888	1.145672	1.041436	
3627	1.590691	1.180300	1.446213	2.164373	2.661269	
2106	0.591345	0.474853	0.370226	0.432887	0.255175	

	Viscera_weight	Shell_weight	Age
3141	-1.583867	-1.644065	-1.543234
3521	-1.565619	-1.626104	-1.387181
883	0.286552	1.538726	1.577830
3627	2.330326	1.377072	0.017298
2106	0.272866	0.906479	1.265723

```
X_Test.head()
```

	Length	Diameter	Height	Whole_weight	Shucked_weight	\
668	0.216591	0.172519	0.370226	0.181016	-0.368878	
1580	-0.199803	-0.079426	-0.466653	-0.433875	-0.443224	
3784	0.799543	0.726798	0.370226	0.870348	0.755318	
463	-2.531611	-2.447709	-2.020857	-1.579022	-1.522362	
2615	1.007740	0.928354	0.848442	1.390405	1.415417	

	Viscera_weight	Shell_weight	Age
668	0.569396	0.690940	0.953617
1580	-0.343004	-0.325685	-0.606915
3784	1.764639	0.565209	0.329404
463	-1.538247	-1.572219	-1.543234
2615	1.778325	0.996287	0.641511

```
Y_Train.head(),Y_Test.head()
```

```
(3141    1
 3521    1
 883     2
 3627    2
 2106    2
Name: Sex, dtype: int64, 668     2
1580    1
3784    2
463     1)
```

```
2615    2
Name: Sex, dtype: int64)
```

Build the Model

```
from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators=10,criterion='entropy')
model.fit(X_Train,Y_Train)
```

```
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
y_predict = model.predict(X_Test)
```

```
y_predict_train = model.predict(X_Train)
```

Training and Testing Model

```
#train the model
```

```
from sklearn.metrics import
accuracy_score,confusion_matrix,classification_report
print('Training accuracy: ',accuracy_score(Y_Train,y_predict_train))
```

```
Training accuracy:  0.9823406165818617
```

```
#test the model
```

```
print('Testing accuracy: ',accuracy_score(Y_Test,y_predict))
```

```
Testing accuracy:  0.5394736842105263
```

```
#Measure the performance using metrics
```

```
pd.crosstab(Y_Test,y_predict)
```

```
col_0    0    1    2
Sex
0       115   30  104
1        41  216   34
2       122   54  120
```

```
print(classification_report(Y_Test,y_predict))
```

	precision	recall	f1-score	support
0	0.41	0.46	0.44	249
1	0.72	0.74	0.73	291
2	0.47	0.41	0.43	296
accuracy			0.54	836
macro avg	0.53	0.54	0.53	836

weighted avg	0.54	0.54	0.54	836
--------------	------	------	------	-----